

Project : Reinforcement Learning (MVA)

# Subgoal discovery in Reinforcement Learning

**Yonatan DELORO** and **Juliette RENGOT**


November, 2018 to January, 2019

# ***Plan of the presentation***

1. *Problem definition : the taxi domain*
2. *Presentation of the proposed solutions*
  - a. Standard Q-learning algorithm on a flat MDP formulation
  - b. Macro Q-Learning with Diverse-Density-based options
    - i. Theoretical explanation
    - ii. Implementation choices
    - iii. Results and performances analysis
  - c. HEXQ algorithm
    - i. Theoretical explanation
    - ii. Implementation choices
    - iii. Results and performances analysis
3. *Conclusion and discussion*

# ***Problem definition : the taxi domain***

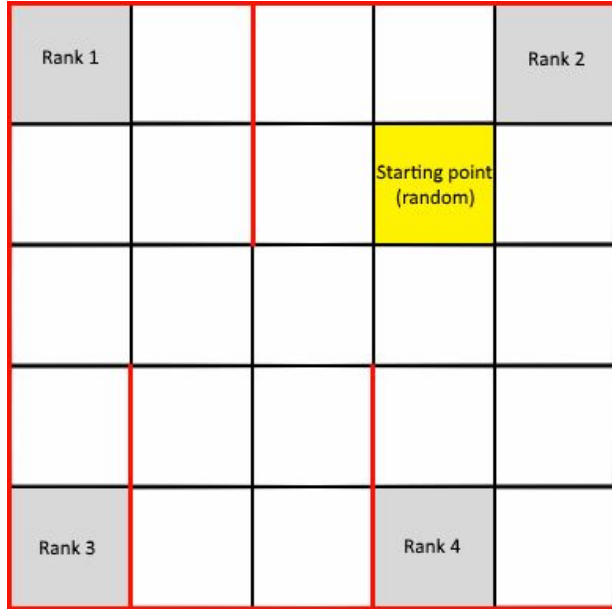
Original Version (Dietterich)



HEXQ test


# ***Problem definition : the taxi domain***

## Original Version (Dietterich)



HEXQ test

## Simplified Version



DDMQ test

(“flat Q-learning”)

***Standard Q-learning  
algorithm on a flat MDP  
formulation***

# ***The flat Q-learning solution***

## **1° Formulating the Taxi Problem as a flat MDP**

*Space of states:*

A state is defined by the triplet :

- the location of the cab in the grid
- the location of the passengers (in the cab / at which rank)
- the location of the destination (at which rank)

*Space of primitive actions:*

left, right, down, up pick up, put down

## **2° Applying standard Q-learning algorithm to solve this MDP**

(“DDMQ”)

# ***Macro Q-Learning with Diverse-Density based options***

Amy McGovern and Andrew G. Barto. “Automatic Discovery Of Subgoals In Reinforcement Learning using Diverse Density”. In: Computer Science Department Faculty Publication Series. 8. (2001).

# ***Theoretical aspects :***

## ***The big picture***

### **Ideas :**

- 1) Mine the set of saved trajectories to :
  - define subgoal candidates as commonalities across successful saved trajectories (“bottlenecks”)
  - build policies aimed at reaching these subgoals
- 2) Enable the agent to follow and assess the built-in policies available from state, instead of primitive actions






# ***Theoretical aspects :***

## ***A. How to discover subgoals ?***



### **1. Looking for bottlenecks is a multiple instance learning problem**

#### *The MIL problem*

##### Positive bags



##### Negative bags



*Goal : find the target concept* ☐

# ***Theoretical aspects :***

## ***A. How to discover subgoals ?***

### **1. Looking for bottlenecks is a multiple instance learning problem**

-A trajectory is modelled as a bag of instances, where an instance typically corresponds to a state visited by the agent

-Successful trajectories = positive bags

-Unsuccessful trajectories = negative bags (with no positive instance by assumption)

-Target concepts = bottleneck region (type of instance responsible for the positivity of the bag it belongs to)

# ***Theoretical aspects :***

## ***A. How to discover subgoals ?***

### **2. Solve the MIL problem using Diverse Density**

Idea: find the region with instances from the most positive bags and the least negative bags

ie. the concept with the highest diverse density as defined below

$$d(t) := Pr(t|B_1^+, ..., B_n^+, B_1^-, ..., B_m^-) = \prod_{i=1}^n Pr(t|B_i^+) \prod_{i=1}^m Pr(t|B_i^-)$$

$$Pr(t|B_i^-) = \prod_{1 \leq j \leq p} (1 - Pr(B_{ij}^- \in c_t))$$

$$Pr(t|B_i^+) = 1 - \prod_{1 \leq j \leq p} (1 - Pr(B_{ij}^+ \in c_t))$$

# ***Theoretical aspects :***

## ***A. How to discover subgoals ?***

### **3. Keep only relevant diverse dense concepts**

- 1) Exclude concepts near terminal state of successful trajectories
- 2) Keep only DD peaks appearing in the first stages of learning and persisting throughout learning

At the end of a trajectory, the agent creates a new bag and looks for the DD peak. If concept  $c$  appears as peak, it updates the running average of how often  $c$  has appeared as a peak  $m_c \leftarrow \lambda(m_c + 1)$

$$m_c > x \frac{\lambda}{1 - \lambda} \quad \Longrightarrow \quad c \text{ is selected as a sub-goal candidate}$$

## ***Theoretical aspects :***

### ***B. Learning with discovered subgoal candidates***

#### **1. Create an option for a subgoal candidate**

An option is a temporally-extended action which the agent executes until a termination condition is satisfied.

It is defined by:

$I$  : the option's input set of states

$\pi$  : the option's own policy

$\beta$  : the probability of termination in each state

# ***Theoretical aspects :***

## ***B. Learning with discovered subgoal candidates***

### **1. Create an option for a subgoal candidate**

We will use options to model built-in policies intended to reach subgoal candidates (target concepts)

In this purpose,

- $I$  : contains the set of states visited on a trajectory during a time window before reaching the subgoal
- $\pi$  : is computed based on a value function learnt using experience replay and a relevant option's reward
- $\beta$  : is 1 if the subgoal is reached or we left the input set

# ***Theoretical aspects :***

## ***B. Learning with discovered subgoal candidates***

### **2. Use Macro Q-Learning to learn both with primitive actions and options**

Primitive action played at t

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left( r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right)$$

Option initiated at t and terminated at t+n

$$Q(s_t, m_t) \leftarrow Q(s_t, m_t) + \alpha \left( \sum_{i=1}^n r_{t+i} \gamma^{i-1} + \gamma^n \max_{a'} Q(s_{t+n}, a') - Q(s_t, m_t) \right)$$

# ***Implementation choices***

## Considered concepts and subgoal candidates

- A concept is defined as reaching an individual state  $(x, y, p)$
- $Pr(B_{ij} \in c_t) = \exp - ||B_{ij} - c_t||^2$
- DD values of concepts computed using positive bags only
- Concepts near the final goal terminal state excluded
- Maximum of one option created per run



# Implementation choices

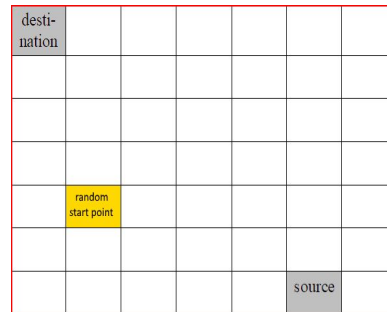
## Considered metric for the probability to belong to a concept

1) First “agnostic” metric (used for report’s results) :

$$||(x, y, p) - (x', y', p')||^2 = a_0(x - x')^2 + a_1(y - y')^2 + a_2(p - p')^2$$

- Problem which still remains even with different scaling factors :

$$||(x, y, 1) - (x, y, 0)||^2 = cste$$



# Implementation choices

## Considered metric for the probability to belong to a concept

1) First “agnostic” metric (used for report’s results) :

$$||(x, y, p) - (x', y', p')||^2 = a_0(x - x')^2 + a_1(y - y')^2 + a_2(p - p')^2$$

- Problem which still remains even with different scaling factors :

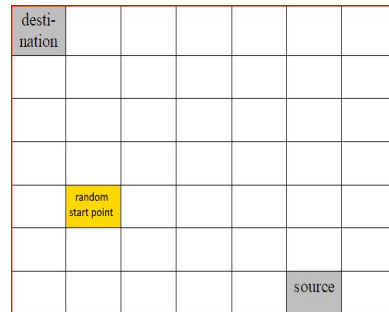
$$||(x, y, 1) - (x, y, 0)||^2 = cste$$

What we would prefer to find as distance instead :

$$||(x, y, 0) - (x', y', 0)||^2 = ||(x', y') - (x, y)||_2^2$$

$$||(x, y, 1) - (x', y', 1)||^2 = ||(x', y') - (x, y)||_2^2$$


$$||(x, y, 0) - (x', y', 1)||^2 = ||(x, y) - (x_{source}, y_{source})||_2^2 + ||(x', y') - (x_{source}, y_{source})||_2^2 + 1$$



# Implementation choices

## Considered metric for the probability to belong to a concept

What we would prefer to find as distance instead :



$$||(x, y, 0) - (x', y', 0)||^2 = ||(x', y') - (x, y)||_2^2$$

$$||(x, y, 1) - (x', y', 1)||^2 = ||(x', y') - (x, y)||_2^2$$

$$||(x, y, 0) - (x', y', 1)||^2 = ||(x, y) - (x_{source}, y_{source})||_2^2 + ||(x', y') - (x_{source}, y_{source})||_2^2 + 1$$

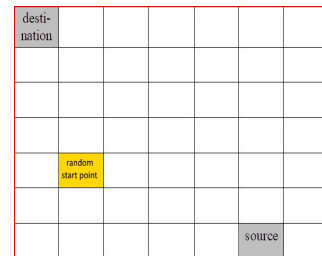
2) Second metric, changing the state representation (“task relevant”)

$$s = \begin{bmatrix} x \\ y \\ p \end{bmatrix} \Rightarrow \begin{bmatrix} p & \times & (x - x_{source}) \\ (1 - p) & \times & (x - x_{source}) \\ p & \times & (y - y_{source}) \\ (1 - p) & \times & (y - y_{source}) \end{bmatrix} := \phi(s) \quad ||s - s'||^2 = ||\phi(s) - \phi(s')||_2^2$$

# Implementation choices

## Considered metric for the probability to belong to a concept

What we would prefer to find as distance instead :



$$||(x, y, 0) - (x', y', 0)||^2 = ||(x', y') - (x, y)||_2^2$$

$$||(x, y, 1) - (x', y', 1)||^2 = ||(x', y') - (x, y)||_2^2$$

$$||(x, y, 0) - (x', y', 1)||^2 = ||(x, y) - (x_{source}, y_{source})||_2^2 + ||(x', y') - (x_{source}, y_{source})||_2^2 + 1$$

2) Second metric, changing the state representation (“task relevant”)

$$s = \begin{bmatrix} x \\ y \\ p \end{bmatrix} \Rightarrow \begin{bmatrix} p & \times & (x - x_{source}) \\ (1 - p) & \times & (x - x_{source}) \\ p & \times & (y - y_{source}) \\ (1 - p) & \times & (y - y_{source}) \end{bmatrix} := \phi(s) \quad ||s - s'||^2 = ||\phi(s) - \phi(s')||_2^2$$

We could retrieve these four features in :

$$P_2 \left( \begin{bmatrix} (x - x_{source}) \\ (y - y_{source}) \\ p \\ 1 - p \end{bmatrix} \right) := \phi(s)$$

$$||s - s'||^2 = \sum_k a_k^2 (\phi(s)_k - \phi(s')_k)^2$$

# ***Implementation choices***

## Computing the option's policy

- Option's reward : 10 is subgoal reached, -1 otherwise (-1 if we leave the input set or if we do not still reach the subgoal)
- Option's value function computed with Q-learning algorithm using option's reward and experience replay over parts of saved trajectories (time window before reaching subgoal)

# ***Results - experimental setup***

## **For both Q-learning and DDMQ :**

- *Environment* : discount factor of 0.95
- *Trajectories* :  $N = 500$ ,  $T_{\max} = 1000$
- *Policy throughout learning* : eps-greedy policy  
eps\_init = 0.4, increase of 0.1 each 100 iters until 0.9
- *Learning rate* : adaptive  
(update of  $Q(s,a)$  : inverse of the number of visits of  $(s,a)$ )

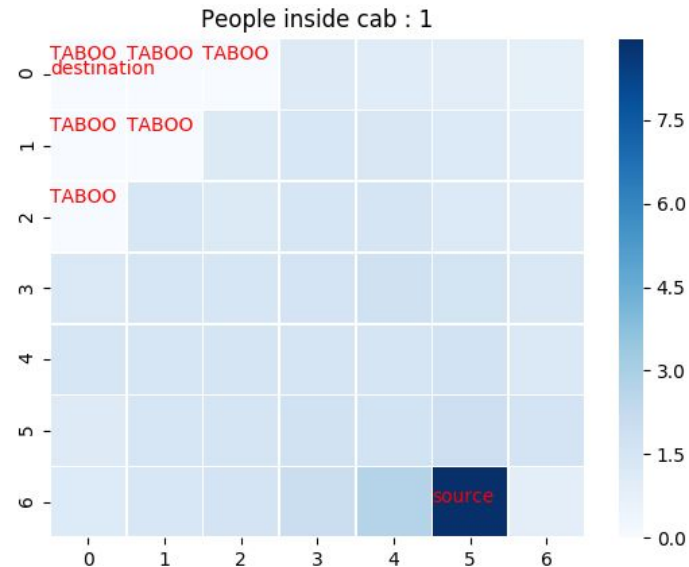
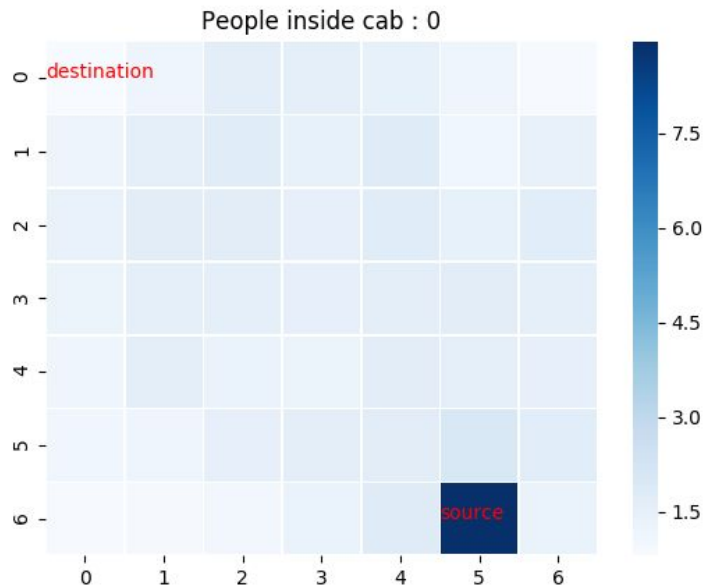
## **DDMQ hyperparameters :**

- Time window of size 5 before reaching subgoal considered for computing option's input set and policy
- Discount factor for option reward = 0.95
- Peak concept frequency parameters :  $(\lambda, x) = (0.9 ; 0.9999)$

$$m_c \leftarrow \lambda(m_c + 1) \quad m_c > x \frac{\lambda}{1 - \lambda}$$

# Results - DDMQ


“agnostic” metric




(a) Running average  $m_c$  of how often each concept  $c$  (taxi location, people presence in the cab) appeared as a peak at episode 50. Average over 30 runs of DDMQ

# Results - DDMQ vs. basic Q-learning

“agnostic” metric



(b) For each algorithm, at a given episode, we plot the average of the discounted cumulative rewards across all previous episodes.




(c) Number of steps to reach the goal throughout the learning. Curbs are average over 30 runs for each of the two algorithms, and moving average over 20 episodes was then applied.




# Results - DDMQ vs. basic Q-learning

“task-relevant” metric



(b) For each algorithm, at a given episode, we plot the average of the discounted cumulative rewards across all previous episodes.



(c) Number of steps to reach the goal throughout the learning. Curbs are average over 30 runs for each of the two algorithms, and moving average over 20 episodes was then applied.

# ***HEXQ algorithm***

**Bernhard Hengst. “Discovering Hierarchy in Reinforcement Learning with HEXQ”.  
In: Proceedings of the Nineteenth International Conference on Machine Learning  
(July 2002), pp. 243–250.**

# ***Theoretical aspects :***

## ***Automatic hierarchical decomposition***

Idea : Decompose a Markov Decision Process by dividing the state space into nested sub-MDP regions.

Conditions :

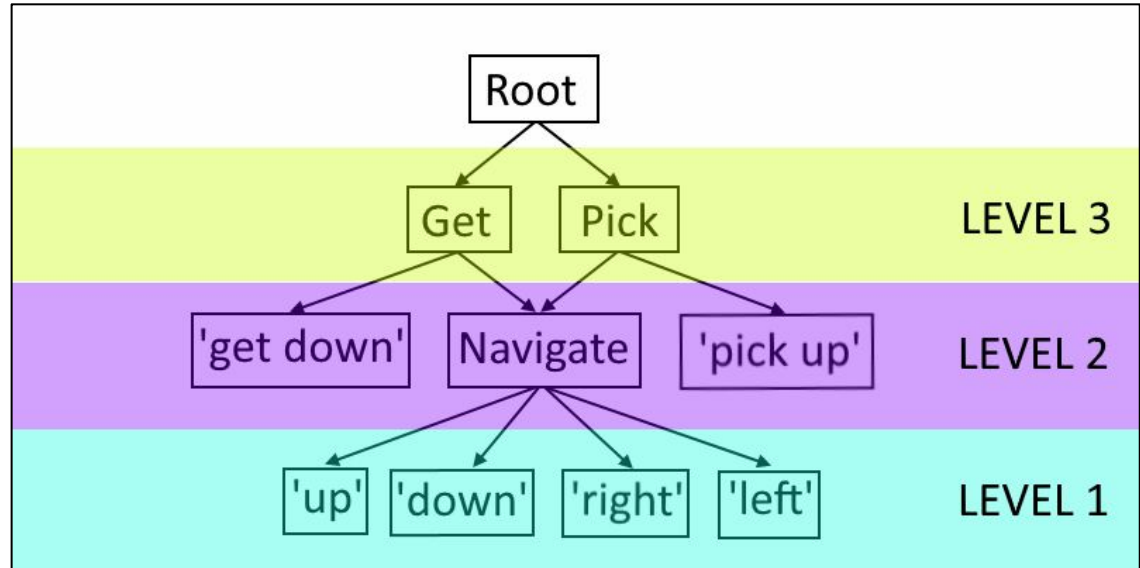
- variables change  $\pm$  frequently
- transition properties retain
- interface between regions can be controlled

# ***Theoretical aspects :***

## ***Automatic hierarchical decomposition***

### Taxi problem

- 3 state variables :
- destination rank  
 $\in \llbracket 0, 3 \rrbracket$
  - passenger source  
 $\in \llbracket 0, 4 \rrbracket$
  - taxi location  
 $\in \llbracket 0, 24 \rrbracket$



# ***Theoretical aspects :***

## ***State abstraction***

$$s^{e+1} = |r^e| x^j + r^e$$

with  $s^{e+1}$  = abstract state value at level  $e+1$

$|r^e|$  = number of regions at level  $e$

$r^e$  = region label from level  $e$

$x^j$  = next most frequent state variable value

# ***Theoretical aspects :***

## ***Action abstraction***

The abstract actions available in each of these abstract states are the policies leading to region exits at the level below.

### Taxi problem :

- 8 exits from bottom level

$(s^1=0, a=\text{pick up}), (s^1=0, a=\text{put down})$   
 $(s^1=4, a=\text{pick up}), (s^1=4, a=\text{put down})$   
 $(s^1=20, a=\text{pickup}), (s^1=20, a=\text{put down})$   
 $(s^1=23, a=\text{pickup}), (s^1=23, a=\text{put down})$

- 4 exits from next level :

$(s^2=0, (s^1 = 0, a=\text{put down}))$   
 $(s^2=4, (s^1 = 0, a=\text{put down}))$   
 $(s^2=0, (s^1 = 20, a=\text{put down}))$   
 $(s^2=0, (s^1 = 23, a=\text{put down}))$

# ***Theoretical aspects :***

## ***The Value Function***

$$Q_{em}^*(s^e, a) = \sum_{s'} T_{se s'}^a (R_{se}^a + V_{em}^*(s'))$$

with  $T_{se s'}^a = \Pr(s'|s^e, a)$

$R_{se}^a = E(\text{primitive reward after } a|s^e, a)$

$s'$  the hierarchical next state

$$V_{em}^*(s) = \max_a (V_{e-1, m_{\{e-1\}}(a)}^*(s) + Q_{em}^*(s^e, a))$$

with  $m_{\{e-1\}}(a)$  the sub-MDP implementing action  $a$

# ***Implementation choices***

## ***Hypothesis of simplification***

We suppose that the hierarchy is already constructed.

In theory, at the beginning of the training, HEXQ should order the variables and find exits at the first level before it can improve its performances over the taxi domain. In the article, this work takes 41 episodes.




# ***Results - experimental setup***

## Parameters :

- Number of episode :  $N = 6000$
- Time horizon :  $T_{\max} = 1000$
- Exploration - exploitation ratio :  $\epsilon_{\text{initial}} = 0.1$   
→ It increases of 0.1 each 200 episodes until 0.9
- *Learning rate* : adaptive  
→ It is based on the inverse of the number of visits of (state, action)

# Results - HEXQ policies




(a) Policy at level 0 when the objective is 0 and people are not in the taxi. (b) Policy at level 0 when the objective is 4 and people are not in the taxi. (c) Policy at level 0 when the objective is 20 and people are in the taxi. (d) Policy at level 0 when the objective is 23 and people are not in the taxi.



(e) Policy at level 1 when destination is 0.



(f) Policy at level 1 when destination is 23.



(g) Policy at level 2.

# Results

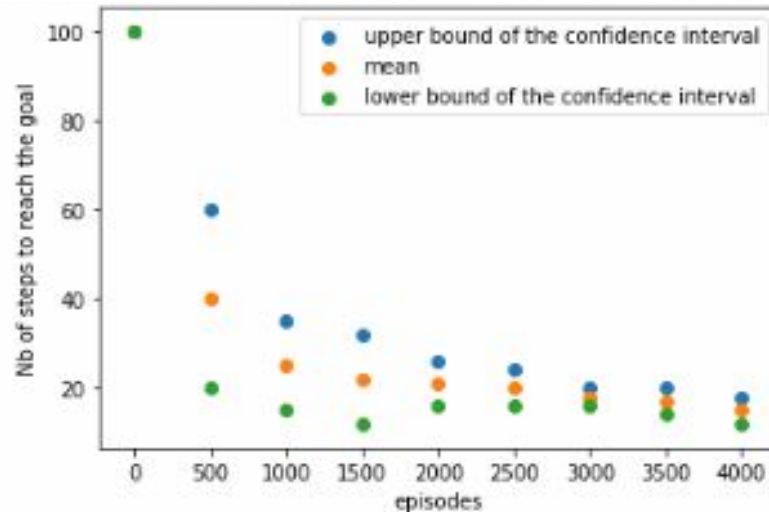
## *HEXQ vs. flat Q-learning*



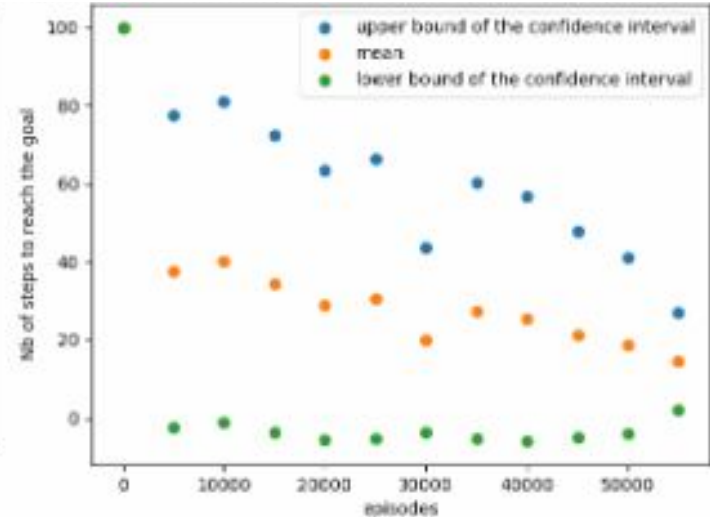
(a) For each algorithm, at a given episode, we plot the average of the discounted cumulative rewards across all previous episodes. (b) Number of steps to reach the goal throughout the learning, for one run of HEXQ algorithm (moving average over 10 episodes was applied).

# Results

## *HEXQ vs. flat Q-learning*

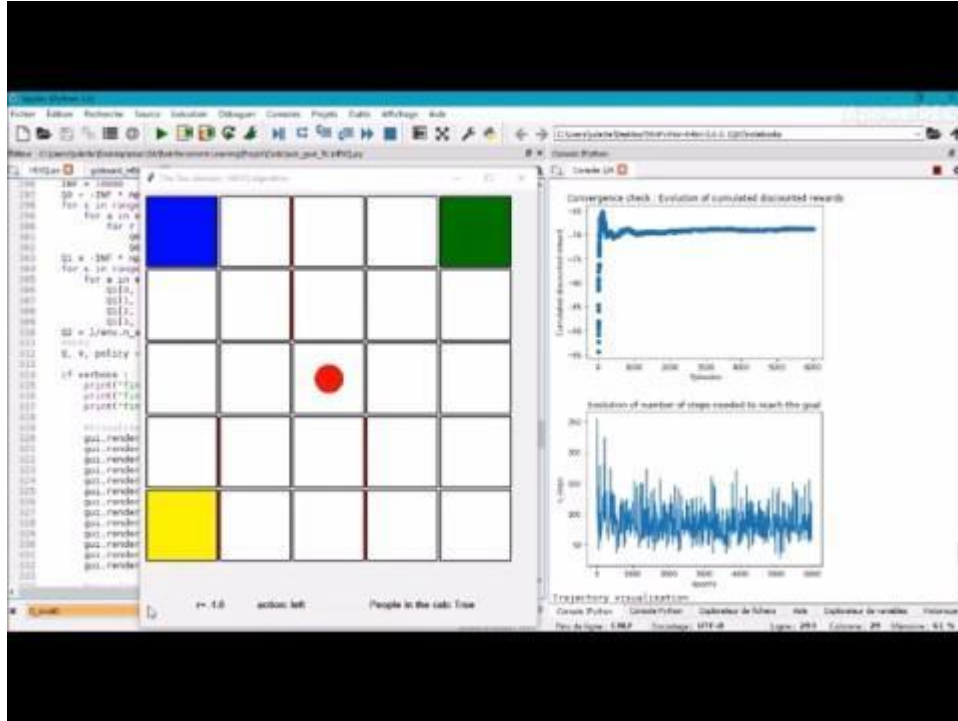


(c) HEXQ: Statistics for 100 trajectories, exploiting the learnt policy, of the number of steps to reach the goal. 68% confidence interval (+/- 1 number of steps to reach the goal. 68% confstd).



(d) Flat Q-learning: Statistics for 100 trajectories, exploiting the learnt policy, of the number of steps to reach the goal. 68% confidence interval (+/- 1 std).

# Results - HEXQ

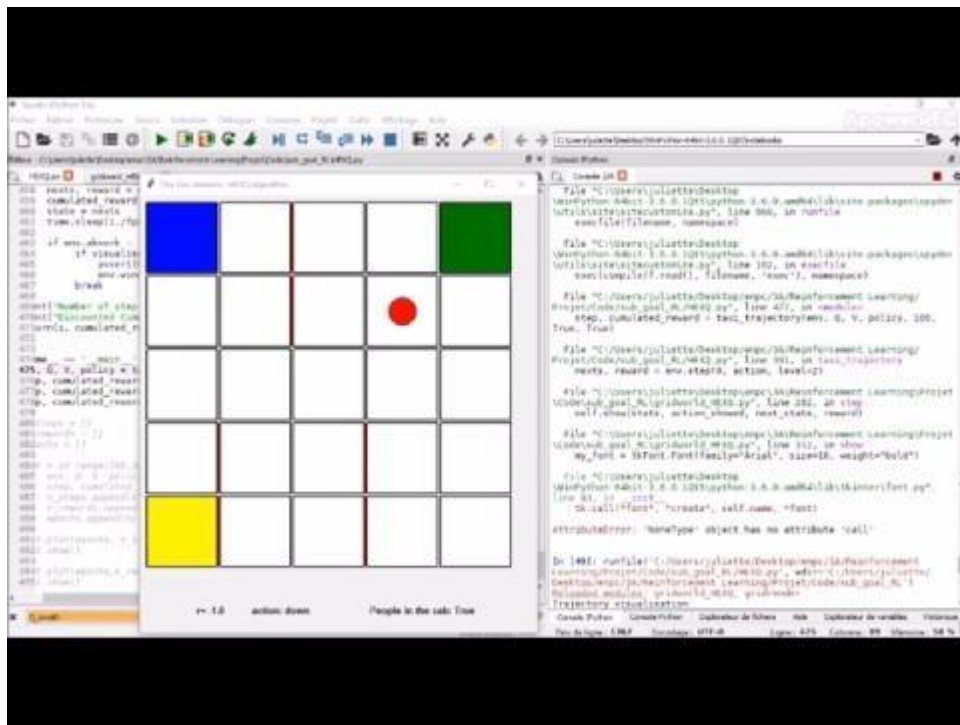


## Demo 1:

Starting point = 20  
passenger source = 4  
destination rank = 0  
Proba\_success = 1

**We reach the goal  
in 17 steps  
with a reward of -12.61.**

# Results - HEXQ



Demo 2:  
Proba\_success = 0.6

**We reach the goal  
in 26 steps  
with a reward of -34.53.**

# ***Conclusion and discussion***

## DDMQ

Interesting when clear bottlenecks in the observation space

Limits :

1° assumes we can reach the main goal using only primitive actions first

2° challenging to define a relevant metric between instances

3° memory cost + number of hyperparameters

## HEXQ

Interesting when repetition of sub-structures in the task

Limits :

1° assumes we can define some variable that vary at longer timescale than others

2° preliminary cost to build the hierarchy (variables frequency analysis and exits discovery)

***Thank you for your attention  
and for the proposal of this project !***



# References

- Andrew H. Fagg Amy McGovern Richard S. Sutton. “Roles of Macro-Actions in Accelerating Reinforcement Learning”. In: In Grace Hopper Celebration of Women in Computing (1997).
- Oded Maron and Tomás Lozano-Pérez. “A Framework for Multiple-Instance Learning”. In: (1998).
- Thomas G. Dietterich. “Hierarchical reinforcement learning with the MAXQ value function decomposition.” In: Journal of Artificial Intelligence Research 13 (2000), pp. 227–303.
- Amy McGovern and Andrew G. Barto. “Automatic Discovery of Subgoals in Reinforcement Learning using Diverse Density”. In: Computer Science Department Faculty Publication Series. 8. (2001).
- Bernhard Hengst. “Discovering Hierarchy in Reinforcement Learning with HEXQ”. In: Proceedings of the Nineteenth International Conference on Machine Learning (July 2002), pp. 243–250.

# ***What we have learnt***

There are many ways we can formulate the same RL task,  
with very variable impacts on the performances

- We should think carefully to the type of actions we define :
  - > *only primitive actions (flat Q-learning)*
  - > *extended in time (DDMQ)*
  - > *at different abstraction levels (HEXQ)*
- How we represent a state can be also important  
(cf. choice of metric in DDMQ)

# Pseudo-code of DDMMQ

## PSEUDO-CODE OF MACRO Q-LEARNING WITH DIVERSE-DENSITY BASED OPTIONS:

### INPUT

$\epsilon$  : exploration/exploitation trade-off schedule

N : number of episodes

Tmax : time horizon of an episode

n : number of steps before reaching a concept within a trajectory, considered for initializing the option input set

lambda : factor inferior to 1 such that 1-lambda is the preference to find a peak concept early

threshold\_score\_concept : threshold on the the running average of how often a concept appears as a peak

gamma\_option : discount factor to compute the policy of an option

### PROCESS

#initialization

$Q \leftarrow$  random matrix of size (49\*2, 6) such that  $Q[s, a] = -\infty$  if the action action is not possible from the state s

Options  $\leftarrow$  [ ]

memory  $\leftarrow$  [ ] #database of trajectories

score\_concepts  $\leftarrow$  vector of zeros of size the number of concepts storing the running average of how often a concept appears as a peak

for k in range(N):

    state  $\leftarrow$  new initial state

    trajectory  $\leftarrow []$

**#MACRO Q-LEARNING**

    while (t < T<sub>max</sub>) and (env is not solved):

        action  $\leftarrow$  choose an action (primitive/option), randomly with a probability  $\epsilon$ , otherwise the one  
        play action and update Q[state\_0, action]

    add trajectory to memory

**#DIVERSE-DENSITY BASED SUBGOAL DISCOVERY TO CREATE NEW OPTION**

    create positive bag from trajectory if successful

    compute the diverse density from all the saved positive bags, and search for diverse density peaks

    for each peak concept c found:

        score\_concepts[c]  $\leftarrow$  lambda \* (score\_concepts[c]+1)

        if (score\_concepts[c] > threshold\_score\_concept) and (c not in taboo states):

            mark concept as subgoal candidate

    select at random one concept c among the subgoal candidates #max one option created per episode

    create a new option o = <I,  $\pi$ ,  $\beta$ > of reaching concept c, and add it to Options

    add a column Q<sub>c</sub> to the Q matrix for the newly created option, such that Q<sub>c</sub>[s] = - $\infty$  if the state s is not in the option input set

$\pi \leftarrow$  compute policy from Q matrix

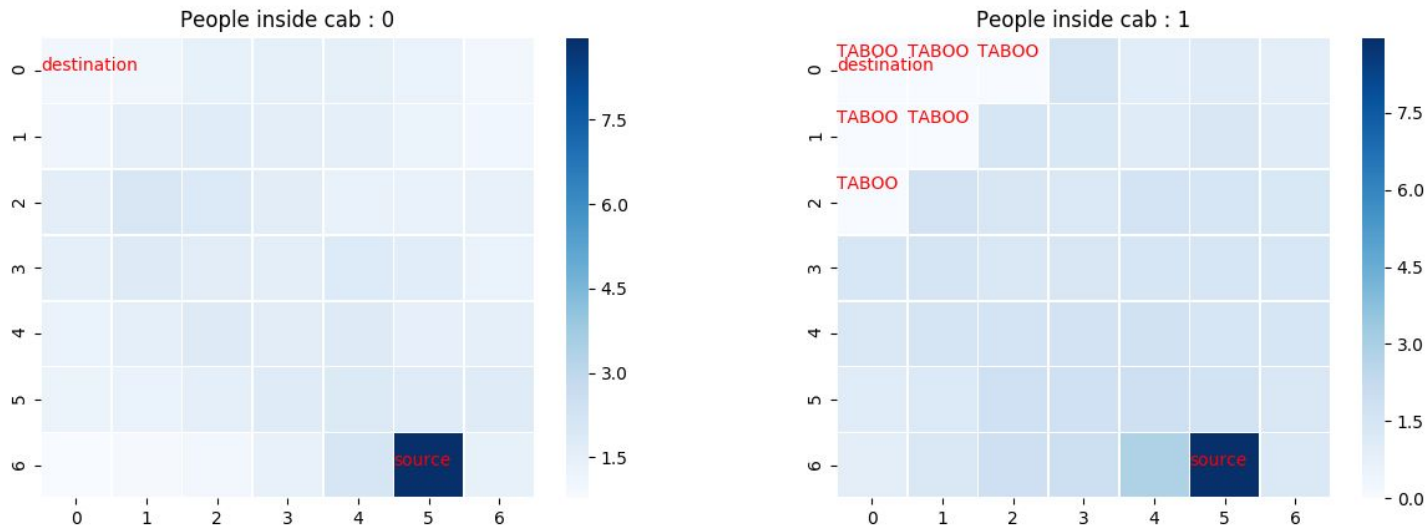
## OUTPUT

final policy :  $\pi$  #associate the index of the primitive action or option to play from each possible state

list of options : Options: #definitions of options (input set, option policy, terminal condition)

# Results - DDMQ

“task-relevant” metric



(a) Running average  $m_c$  of how often each concept  $c$  (taxi location, people presence in the cab) appeared as a peak at episode 50. Average over 30 runs of DDMQ

# *Implementation choices*

## *Pseudo-code of HEXQ training*

### **INPUT**

eps = Exploration/exploitation factor

N = Number of epochs

Tmax = Time horizon

### **PROCESS**

#initialisation

env  $\leftarrow$  Environment of the taxi problem

Q\_level0  $\leftarrow$  random matrice of size (2, 4, 25, 6) such as  $\forall(i, j, k, l), Q\_level0[i, j, k, l] = -\infty$  if the action  $l$  is not possible from the state  $k$

Q\_level1  $\leftarrow$  random matrice of size (4, 5, 8) such as  $\forall(i, j, k), Q\_level1[i, j, k] = -\infty$  if the action  $k$  is not possible from the state  $j$

Q\_level2  $\leftarrow$  random matrice of size (4, 4) such as  $\forall(i, j), Q\_level2[i, j] = -\infty$  if the action  $j$  is not possible from the state  $i$



# *Implementation choices*

## *Pseudo-code of HEXQ training*

```
for k in range(N):  
    New initialisation of env  
    while ((t < Tmax) and env is not solved):  
        people ← boolean storing True if people are inside the taxi and False otherwise  
        objective ← destination_rank if people are inside the taxi and source_rank otherwise  
        state_level0 ← taxi location  
        if current hierarchical_level is 0 :  
            action ← choose a primitive action, randomly with a probability  $\epsilon$  or the one that maximise  
                Q_level0[people, objective, state_level0, :]  
            apply the chosen action  
            Q_level0[people, objective, :, :] ← Updating with Qlearning  
            V_level0[people, objective, :] ← Value function associated to Q_level0  
  
            if we reach an exit state of level 1 :  
                Q_level1[people, objective, :, :] ← Updating, from Q_level0 and V_level0, with  
                    hierarchical formulas  
            if we are in exit state of level 2 :  
                Q_level2 ← Updating, from Q_level1 and V_level1, with hierarchical formulas
```

# *Implementation choices*

## *Pseudo-code of HEXQ training*

```
if current hierarchical_level is 1 :  
    state_level1  $\leftarrow$  people location  
    action  $\leftarrow$  choose an abstract action of level 1, randomly with a probability  $\epsilon$  or the one that  
    maximise  $Q\_level1[objective, state\_level1, :]$   
    Apply the chosen action  
     $Q\_level1[objective, :, :] \leftarrow$  Updating with Qlearning  
     $V\_level1[objective, :] \leftarrow$  Value function associated to  $Q\_level1$   
  
    if we are in exit state of level 2 :  
         $Q\_level2 \leftarrow$  Updating, from  $Q\_level1$  and  $V\_level1$ , with hierarchical formulas  
  
if current hierarchical_level is 2 :  
    state_level2  $\leftarrow$  destination location  
    action  $\leftarrow$  choose an abstract action of level 2, randomly with a probability  $\epsilon$  or the one that  
    maximise  $Q\_level2[state\_level2, :]$   
    Apply the chosen action  
     $Q\_level2 \leftarrow$  Updating with Qlearning  
     $V\_level2 \leftarrow$  Value function associated to  $Q\_level2$   
  
 $t \leftarrow t+1$ 
```

### OUTPUT

Q matrices :  $Q\_level0, Q\_level1, Q\_level2$   
associated value functions :  $V\_level0, V\_level1, V\_level2$   
associated policies:  $policy\_level0, policy\_level1, policy\_level2$



# *Implementation choices*

## *Pseudo-code of HEXQ testing*

### PROCESS

New initialisation

# Looking for the destination and people locations

action  $\leftarrow$  policy[2][index of destination rank]

found\_destination  $\leftarrow$  objective associated to action

action  $\leftarrow$  policy[1][index of found destination rank][1 + index of source rank]

found\_people  $\leftarrow$  objective associated to action

action  $\leftarrow$  policy[1][index of found people rank][0]

found\_destination  $\leftarrow$  objective associated to action

#Navigation

objective  $\leftarrow$  index of found people rank

for i in range(Tmax):

    people  $\leftarrow$  variable storing 1 if people are in the taxi and 0 otherwise

    if people==1 :

        objective  $\leftarrow$  index of found destination rank


    action  $\leftarrow$  policy[0][people][objective][taxi location state]

    apply action


    visualise action

# Results - DDMQ vs. basic Q-learning

“task-relevant” metric  
transition probability : 0.8



(b) For each algorithm, at a given episode, we plot the average of the discounted cumulative rewards across all previous episodes.



(c) Number of steps to reach the goal throughout the learning. Curbs are average over 30 runs for each of the two algorithms, and moving average over 20 episodes was then applied.